



Introduction to Machine Learning (CS419M)

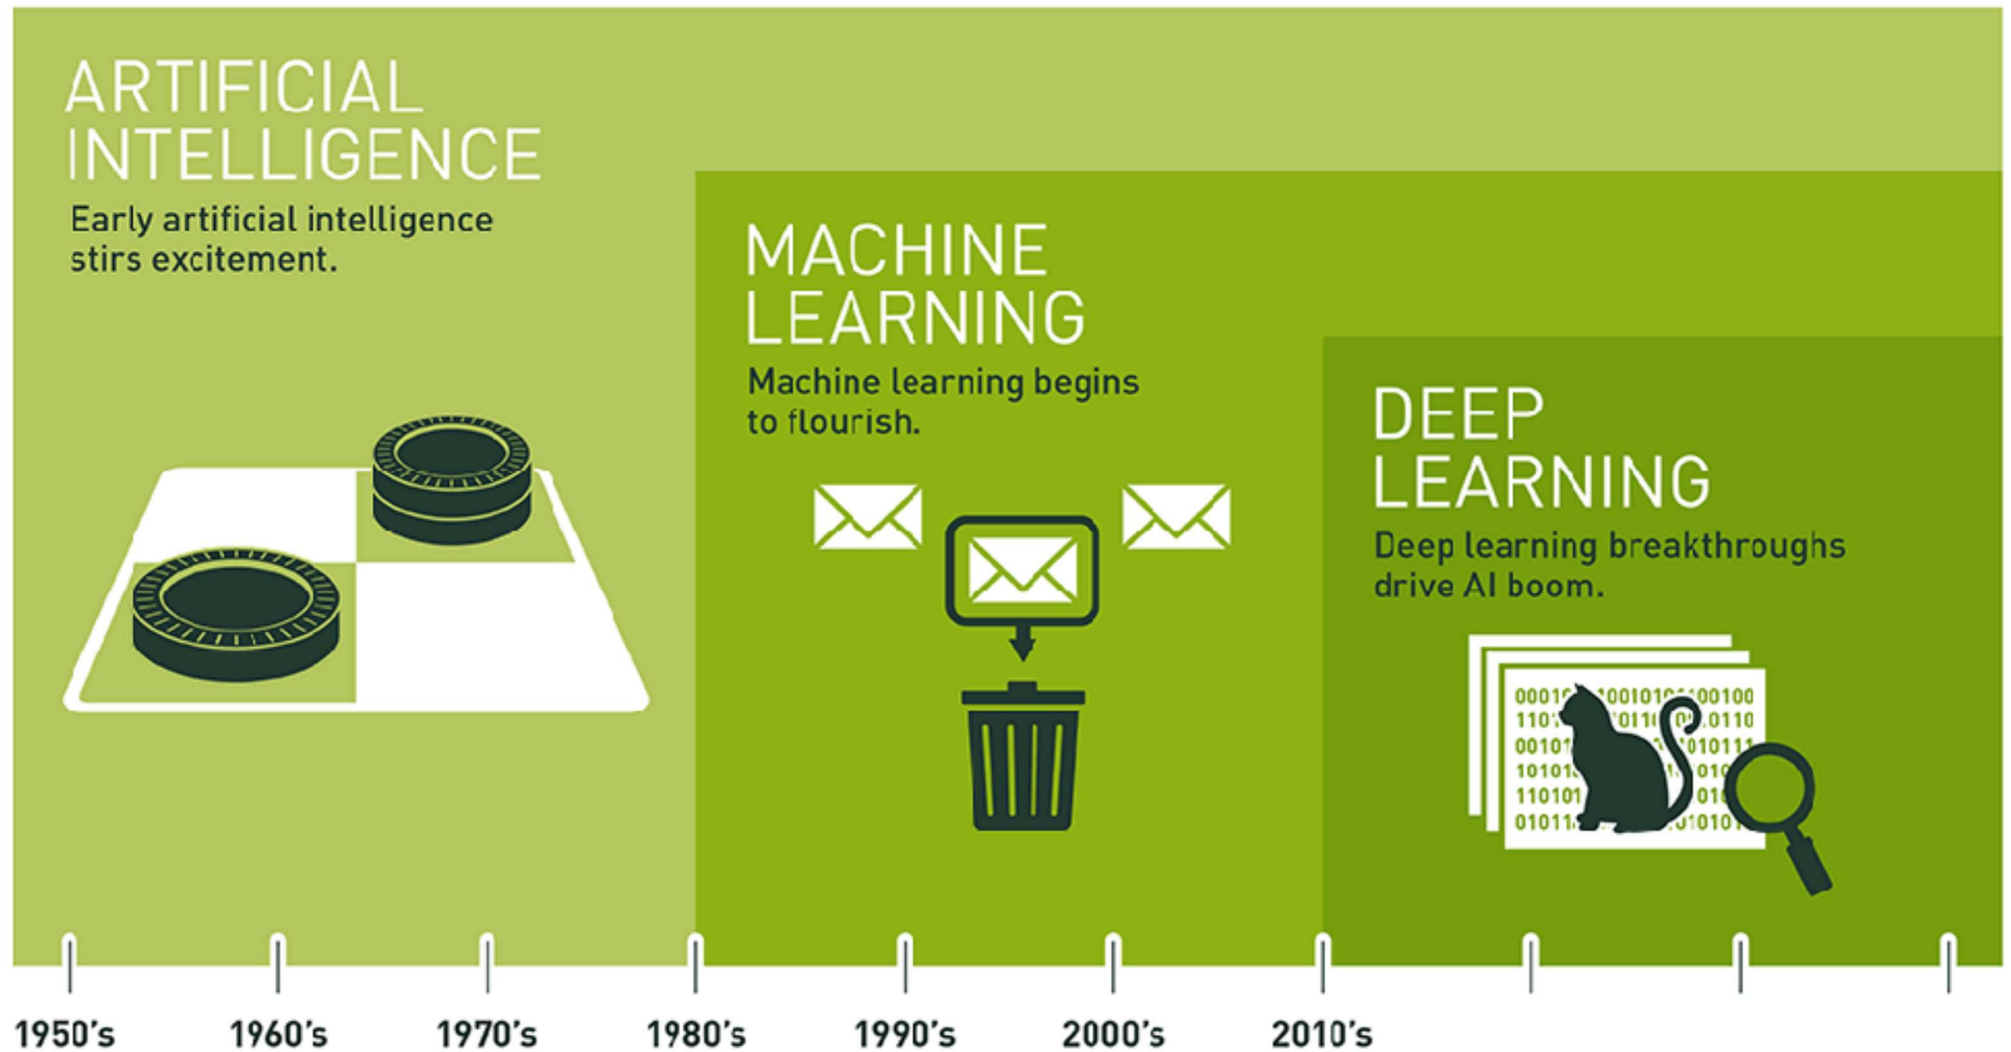
Lecture 1:

- What is learning?
- Supervised vs. unsupervised learning
- Basic course administration and trivia

What is Machine Learning?

- Machine Learning (ML) is a sub-field of computer science that evolved from the study of **pattern recognition** and **computational learning theory** in artificial intelligence.
- Using *algorithms* that iteratively learn from *data*
- Allowing computers to discover *patterns* without being explicitly programmed where to look

Relationship between AI, ML, DL



ML and Statistics?

Glossary

Machine learning

Statistics

network, graphs

model

weights

parameters

learning

fitting

generalization

test set performance

supervised learning

regression/classification

unsupervised learning

density estimation, clustering

large grant = \$1,000,000

large grant = \$50,000

nice place to have a meeting:
Snowbird, Utah, French Alps

nice place to have a meeting:
Las Vegas in August

When do we need ML? (I)

- For tasks that are easily performed by humans but are complex for computer systems to emulate
 - **Vision:** Identify faces in a photograph, objects in a video or still image, etc.
 - **Natural language:** Translate a sentence from Hindi to English, question answering, identify sentiment of text, etc.
 - **Speech:** Recognise spoken words, speaking sentences naturally
 - **Game playing:** Play games like chess, Go, Dota.
 - **Robotics:** Walking, jumping, displaying emotions, etc.
 - Driving a car, navigating a maze, etc.

When do we need ML? (II)

- For tasks that are beyond human capabilities
 - Analysis of large and complex datasets
 - E.g. IBM Watson's Jeopardy-playing machine



Machine Learning

- Ability of computers to “learn” from “data” or “past experience”

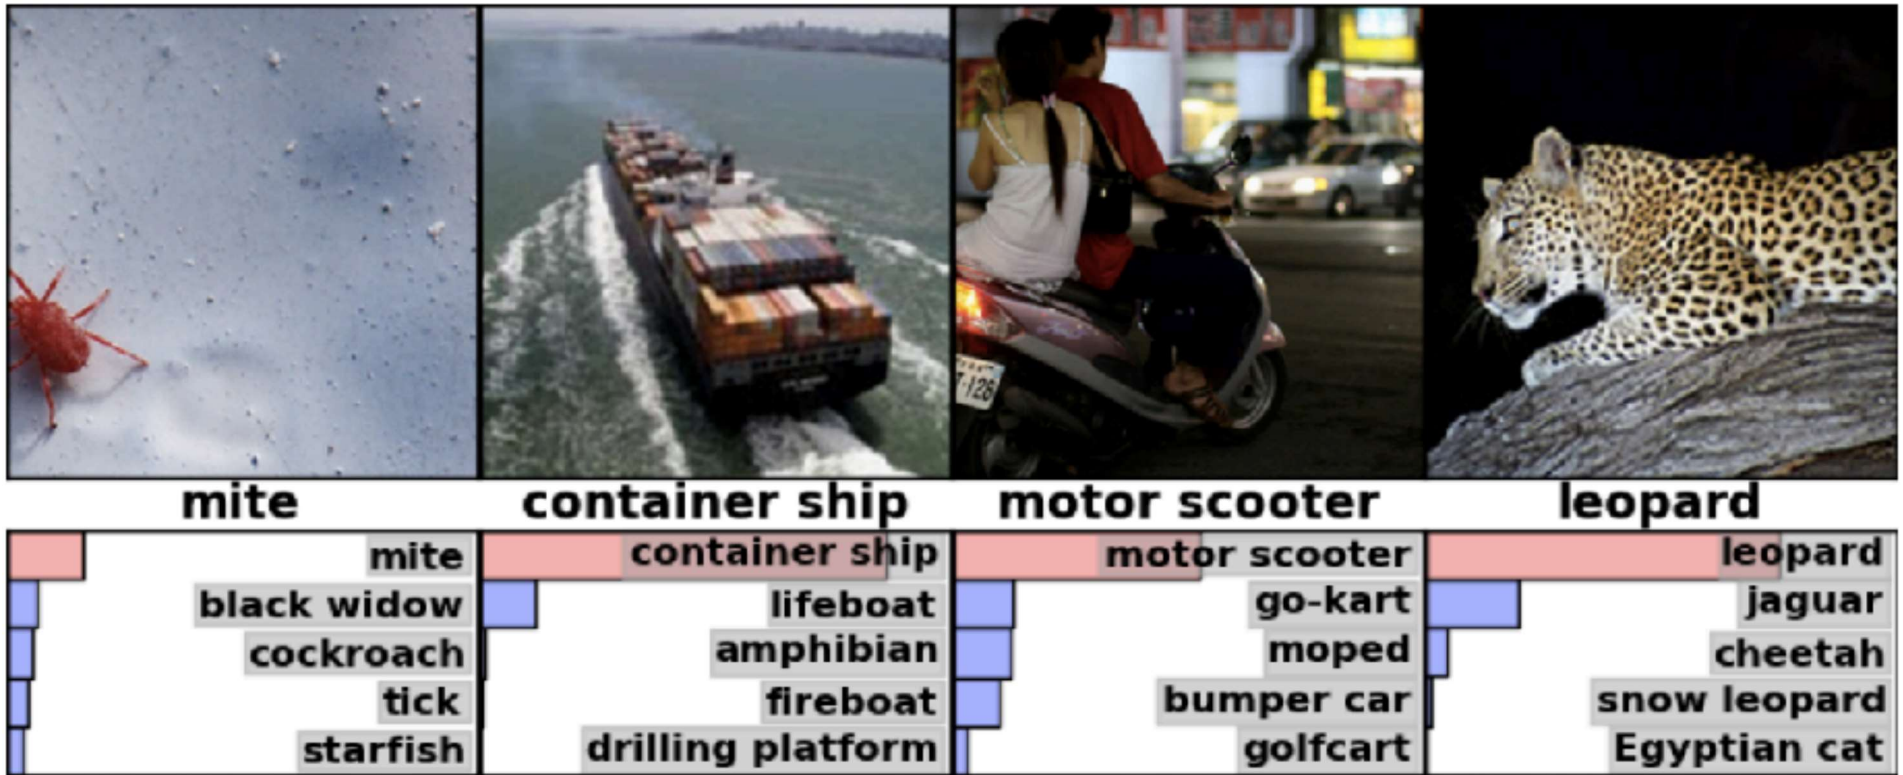
Machine Learning

- Ability of computers to “learn” from “**data**” or “past experience”
- **data**: Comes from various sources such as sensors, domain knowledge, experimental runs, etc.

Machine Learning

- Ability of computers to “**learn**” from “**data**” or “past experience”
- **data**: Comes from various sources such as sensors, domain knowledge, experimental runs, etc.
- **learn**: Make *intelligent* predictions or decisions based on data by optimizing a **model**

Example: Image Recognition



Machine Learning

- Ability of computers to “**learn**” from “**data**” or “past experience”
- **data**: Comes from various sources such as sensors, domain knowledge, experimental runs, etc.
- **learn**: Make intelligent predictions or decisions based on data by optimizing a **model**
 1. Supervised learning: decision trees, neural networks, etc.

Machine Learning

- Ability of computers to “**learn**” from “**data**” or “past experience”
- **data**: Comes from various sources such as sensors, domain knowledge, experimental runs, etc.
- **learn**: Make intelligent predictions or decisions based on data by optimizing a **model**
 1. Supervised learning: decision trees, neural networks, etc.
 2. Unsupervised learning: k-means clustering, etc.

Machine Learning

- Ability of computers to “**learn**” from “**data**” or “past experience”
- **data**: Comes from various sources such as sensors, domain knowledge, experimental runs, etc.
- **learn**: Make intelligent predictions or decisions based on data by optimizing a **model**
 1. Supervised learning: decision trees, neural networks, etc.
 2. Unsupervised learning: k-means clustering, etc.
 3. *Reinforcement learning: Not covered in this course.*

Course Specifics / Administration / Trivia

Prerequisites

No official prerequisites.

Should be comfortable with

- basic probability theory
- linear algebra
- multivariable calculus
- programming (for assignments and project)

Course Webpage

<https://www.cse.iitb.ac.in/~pjyothi/cs419/>

Introduction to Machine Learning, Spring 2020

Course Description

Welcome to "*Introduction to Machine Learning 419(M)*". In this undergraduate-level course, you will be introduced to the foundations of machine learning along with a slew of popular machine learning techniques. This will also give you insights on how to apply machine learning to solve a new problem.

This course is open to any non-CSE undergraduate student who wants to do a minor in CSE. There are no prerequisites.



Course Info

Time: Wednesdays, Fridays, 9.30 am to 10.55 am

Venue: CC 103

Instructor: Preethi Jyothi. You can email me at [pjyothi \[at\] cse \[dot\] iitb \[dot\] ac \[dot\] in](mailto:pjyothi@cse.iitb.ac.in)

TAs (expand [at] cse] as in the instructor's email ID above):

Srijon Sarkar (email: [srijon \[at\] cse](mailto:srijon@cse.iitb.ac.in))

Navya Muttineni (email: [mnavya \[at\] cse](mailto:mnavya@cse.iitb.ac.in))

Shivam Sood (email: [ssood \[at\] cse](mailto:ssood@cse.iitb.ac.in))

Mayur Warialani (email: [mayurwarialani \[at\] cse](mailto:mayurwarialani@cse.iitb.ac.in))

Achari Rakesh Prasanth (email: [rakeshprasanth \[at\] cse](mailto:rakeshprasanth@cse.iitb.ac.in))

Rishabh Kumar (email: [krrishabh \[at\] cse](mailto:krrishabh@cse.iitb.ac.in))

Instructor office hours (in CC 221): 4 to 5 pm on Wednesdays

TA office hours: TBA

Course logistics

Reading: All mandatory reading will be freely available online and posted on the course website.

Textbooks (available online):

1. [Understanding Machine Learning](#). Shai Shalev-Shwartz and Shai Ben-David. Cambridge University Press. 2017.
2. [The Elements of Statistical Learning](#). Trevor Hastie, Robert Tibshirani and Jerome Friedman. Second Edition. 2009.

Attendance: 60% minimum attendance. Counts towards participation points. Strongly advised to attend class. Lot of material will be covered in class, which will not be on the slides.

Personnel and Academic Integrity

Course TAs: Srijon Sarkar, Navya Muttineni, Shivam Sood, Mayur Warialani, Achari Rakesh Prasanth, Rishabh Kumar

Communication:

We will use Moodle for all course-related announcements.

My office hours: 4 pm to 5 pm on Fridays

TA's office hours: TBA. Will be spread out over the week.

Code of conduct:

Abide by an honour code and not be involved in any plagiarism. If caught for copying or plagiarism, name of both parties will be handed over to the Disciplinary Action Committee (DAC)¹.

¹<http://www1.iitb.ac.in/newacadhome/punishments201521July.pdf>

Course Syllabus

Provide an overview of machine learning and well-known ML techniques. We will briefly cover some ML applications as well.

Some Topics:

- Basic foundations of ML, classification/regression, Naive Bayes' classifier, linear and logistic regression
- Supervised learning: Decision trees, perceptron, support vector machines, neural networks.
- Unsupervised learning: k-means clustering, EM algorithm.
- Other topics: feature selection, dimensionality reduction, boosting, bagging.
- Brief introduction to ML applications in computer vision, speech and natural language processing.

Evaluation (subject to minor changes)

Two programming assignments	(20%)
Two quizzes	(20%)
Midsem Exam	(20%)
Final Exam	(25%)
Project	(10%)
Participation	(05%)

Audit requirements:

Both assignments, both quizzes, participation points.

Score 50% or above to successfully audit the course.

Final Project

Team: 2-3 members. Individual projects are highly discouraged.

Project details:

- Apply the techniques you studied in class to any interesting problem of your choice
- Think of a problem early and work on it throughout the course. Project milestones will be posted on Moodle.
- Examples of project ideas: auto-complete code, generate song lyrics, help irctc predict ticket prices, etc.
- Feel free to be creative; consult with TAs/me if it's feasible

Datasets abound...

Kaggle: <https://www.kaggle.com/datasets>

Welcome to Kaggle Datasets

The best place to discover and seamlessly analyze open data




Discover

Use the search box to find open datasets on everything from government, health, and science to popular games and dating trends.



Explore

Execute, share, and comment on code for any open dataset with our in-browser analytics tool, [Kaggle Kernels](#). You can also download datasets in an easy-to-read format.



Create a Dataset

Contribute to the open data movement and connect with other data enthusiasts by clicking "[New Dataset](#)" to publish an open dataset of your own.

[Learn More](#)

[New Dataset](#)

Datasets abound...

Kaggle: <https://www.kaggle.com/datasets>

Another good resource: <http://deeplearning.net/datasets/>

Popular resource for ML beginners:
<http://archive.ics.uci.edu/ml/index.php>

Interesting datasets for computational journalists:
<http://cjlabs.stanford.edu/2015/09/30/lab-launch-and-data-sets/>

Speech and language resources:
www.openslr.org/

... and so do ML libraries/toolkits

scikit-learn, openCV, Keras, Tensorflow, NLTK, etc.

Typical ML approach

- How do we approach an ML problem?
 - **Modeling**: Use a model to represent the task
 - **Decoding/Inference**: Given a model, answer questions with respect to the model
 - **Training**: The model could be parameterized and the parameters are estimated using data

How do we know if our model's any good?

- **Generalization:** Does the trained model produce good predictions on examples beyond the training set?
- We should be careful not to *overfit* the training data
 - Occam's Razor: All other things being equal, pick the simplest solution
- These concepts will be made more precise in later classes